

Muhammad Muhaimin

☎ +1 (647) 608-5061 ✉ muhaiminmemon@gmail.com 💻 [muhaiminmemon](#) 🌐 [muhaiminmemon](#) 📁 [Portfolio](#)

EDUCATION

University of Toronto

Expected Graduation, April 2027

Bachelor's in Computer Science & Management

Toronto, Ontario

- **Summary:** Third-year student with 3+ years of professional experience across 4 engineering roles including Nokia R&D and Autodesk, spanning AI/ML systems, cloud infrastructure, and full-stack development. Canadian citizen.

TECHNICAL SKILLS

Languages: Python, JavaScript/TypeScript, Java, C, C++, Go, SQL, PowerShell, HTML/CSS

AI/ML: LangGraph, LangChain, MCP, Ollama, PyTorch, Hugging Face, ChromaDB, LLM APIs (OpenAI, Anthropic)

Backend/Frameworks: FastAPI, Node.js/Express, React, Next.js, PostgreSQL, Supabase

DevOps/Cloud: Docker, Kubernetes, OpenShift, Helm, AWS, Azure, Git

EXPERIENCE

Machine Learning Operations Engineer

May 2026 – Aug 2026

Autodesk

Toronto, ON

- Building **MLOps** pipelines and infrastructure to deploy and monitor **ML models** and **LLM-based agents** in production

Full-Stack Cloud & Machine Learning Engineer

Jan 2026 – Apr 2026

Nokia Research & Development

Ottawa, ON

- Architected and shipped a **LangGraph** multi-agent AI system for Nokia 5G troubleshooting, cutting query resolution from **2–3 min to 30–40s** via **ChromaDB**, **Redis**, and **Kubernetes** integration over **MCP**
- Identified inference latency bottlenecks and built a two-tier **Redis** semantic cache with cosine-similarity lookup, reducing **LLM** inference costs by **~60%**
- Engineered a **Crossplane**-based edge deployment system with custom **Kubernetes operators** in **Go**, reducing provisioning time from **2 hours to under 5 minutes (~95%)** and eliminating manual configuration errors
- Built a fault-tolerant **Python Kafka** sync service between hub and edge clusters, sustaining **40+** msgs/sec while eliminating replication loops via header-based deduplication and **LRU** caching

Machine Learning Engineer

May 2025 – July 2025

Outamotion

Remote

- Designed and deployed a **computer vision pipeline** integrating **PyMuPDF OCR** with custom **NLP models** to extract structured data from unstructured documents, achieving an **F1-score of 0.95**
- Built **RAG-based retrieval system** with **LlamaIndex** that improved information search accuracy by **40%** over keyword search, enabling faster retrieval of relevant clauses across the repository of documents
- Benchmarked and optimized **LLMs (GPT-4, Llama-3, Mistral-7B)** for use in the RAG pipeline, using **embedding similarity** and response scoring, identifying the most reliable model for production use

Software Engineer

September 2024 – December 2024

DataAnnotation.tech

Toronto, ON

- Developed **Python** scripts to process and validate **AI/ML** training datasets, reducing manual processing time by **30%**
- Fine-tuned **transformer-based LLMs** via **prompt engineering** and **hyperparameter optimization**, boosting accuracy in **sentiment analysis** and classification tasks

IAM Engineer

Jan 2024 – August 2024

Symcor

Mississauga, ON

- Engineered a **PowerShell** automation tool leveraging **Microsoft Graph API** to create, customize, and analyze **400+** Azure access reviews, reducing manual audit effort by **90%**
- Designed and maintained **Python ETL pipelines** via **Graph API** ingesting **120+** feeds and validating identity records across **Azure Cloud**, **GitLab**, and **Active Directory**, flagging **~500** inconsistencies for remediation

PROJECTS

Tessera | Python, OpenAI API, ChromaDB, sentence-transformers, Pydantic, Unsloth

[GitHub](#) 

- Built an open-source synthetic dataset generation library that replaces human annotation; converts a task description into a validated fine-tuning dataset across 4 NLP task types (classification, extraction, instruction-following, RAG/QA)
- Tessera-trained **Llama-3.2-3B** achieves **97.1%** of real-data F1 on **Banking77** for **\$0.40** vs. \$125 crowd-sourced annotation (**300× cheaper**); **98%** hallucination refusal accuracy on enterprise RAG evaluation
- Shipped as a production open-source library with CI/CD, thread-safe **ThreadPoolExecutor** parallel generation, multi-provider LLM routing (OpenAI, Anthropic, Together, Groq), and structured exception hierarchy

Wisp | Javascript, React, Node.js, Express, Supabase, GPT API, Google Calendar API, Docker, Azure

[Demo](#) 

- Built and shipped an AI-powered Chrome extension that blocks distractions in real time with **90%+** accuracy by analyzing browsing activity against user tasks pulled from the **Google Calendar API**, using the **GPT API** with prompt fine-tuning
- Developed a scalable **REST API** in **Node.js/Express** handling **10,000+** requests/day, deployed via **Docker** on **Azure**
- Implemented user authentication via **Google OAuth** and built a dashboard with **React/Tailwind CSS** for real-time task progress and analytics